

RISC-V Privileged Architecture

Muhammad Tahir

Lecture 13

Electrical Engineering Department
University of Engineering and Technology Lahore

Contents

- ① CSR/Special Instructions
- ② Privileged Modes
- ③ Trap/Interrupt CSRs
- ④ Exception Handling
- ⑤ Startup File
- ⑥ Updated Datapath

CSR Instructions: I Type

- Format: I-type




Table 1: I type CSR instructions.

| Instruction | Operation | Type | Illustration |
|-------------|-------------------------------|------|---------------------|
| CSRRW | Atomic Read and Write | I | CSRRW rd, csr, rs1 |
| ✓ CSRRS | Atomic Read and Set bit | I | CSRRS rd, csr, rs1 |
| CSRRC | Atomic Read and Clear bit | I | CSRRC rd, csr, rs1 |
| CSRRWI | Atomic R/W Imm | I | CSRRWI rd, csr, imm |
| CSRRSI | Atomic Read and Set bit Imm | I | CSRRSI rd, csr, imm |
| CSRRCI | Atomic Read and Clear bit Imm | I | CSRRCI rd, csr, imm |

CSR Instructions: I Type Cont'd

- Format: I-type



- $\text{opcode} = \text{csr} \leftarrow \text{rs1}, \quad \text{rd} \leftarrow \text{csr}$
 $\text{func3} = \text{CSRRW}$
- Pseudo-instructions to read and write a CSR
 - $\text{CSRR} \text{ csr, rd (func3 = CSRRS, rs1 = x0): } \text{rd} \leftarrow \text{csr}$
 - $\text{CSRW} \text{ rs1, csr (func3 = CSRRW, rd = x0): } \text{csr} \leftarrow \text{rs1}$

Special Instructions: Environment Calls

- Format: I-type
- Used to generate an **exception** to invoke Operating System (ECALL) and debugger (EBREAK)




Table 2: Environment Call Instructions.

| Instruction | Operation | Type | Illustration |
|-------------|-------------------|------|--------------|
| ECALL | Environment Call | I | ECALL |
| EBREAK | Environment Break | I | EBREAK |

Special Instructions: FENCE

- Format: I-type
- Provides (memory) synchronization barriers for data memory (fence) and instruction memory (FENCE.I)





Table 3: Synch Instructions.

| Instruction | Operation | Type | Illustration |
|-------------|---------------------------------|------|--------------|
| FENCE | Synchronize data accesses | I | fence |
| FENCE.I | Synchronize instruction fetches | I | fence.i |

Special Instructions: Trap-Return Instructions

- Format: I-type
- **URET** returns from an exception handler in user mode
- **SRET** returns from an exception handler in supervisor mode
- **MRET** returns from an exception handler in machine mode



RISC V Privilege Levels

- Privilege levels are used to provide protection between different components of the software stack
- At a given time, a RISC-V hardware thread (hart) is running at some privilege level
- Privilege level is encoded as a mode in one or more CSRs (control status registers)

| Level | Encoding | Name | Abbreviation |
|-------|----------|------------------------|--------------|
| 0 | 00 | User/application level | U |
| 1 | 01 | Supervisor level | S |
| 2 | 10 | Reserved | |
| 3 | 11 | Machine level | M |

Table 4: RISC V privilege levels.

RISC V Privilege Levels Cont'd

- Supported modes at different privilege levels

| No. of Level | Supported Modes | Usage |
|--------------|-----------------|--|
| 1 | M | For simple embedded systems |
| 2 | M, U | For secure embedded systems |
| 3 | M, S, U | For systems running Linux type operating systems |

Table 5: Combinations of privilege levels that can be supported.

CSRs Address Allocation

| CSR Address | | | Hex | Use and Accessibility |
|-----------------|-------|-------|-------------|-----------------------|
| [11:10] | [9:8] | [7:4] | | |
| User CSRs | | | | |
| 00 | 00 | XXXX | 0x000-0x0FF | Standard read/write |
| 01 | 00 | XXXX | 0x400-0x4FF | Standard read/write |
| 10 | 00 | XXXX | 0x800-0x8FF | Custom read/write |
| 11 | 00 | 0XXX | 0xC00-0xC7F | Standard read-only |
| 11 | 00 | 10XX | 0xC80-0xCBF | Standard read-only |
| 11 | 00 | 11XX | 0xCC0-0xCFF | Custom read-only |
| Supervisor CSRs | | | | |
| 00 | 01 | XXXX | 0x100-0x1FF | Standard read/write |
| 01 | 01 | 0XXX | 0x500-0x57F | Standard read/write |
| 01 | 01 | 10XX | 0x580-0x5BF | Standard read/write |
| 01 | 01 | 11XX | 0x5C0-0x5FF | Custom read/write |
| 10 | 01 | 0XXX | 0x900-0x97F | Standard read/write |
| 10 | 01 | 10XX | 0x980-0x9BF | Standard read/write |
| 10 | 01 | 11XX | 0x9C0-0x9FF | Custom read/write |
| 11 | 01 | 0XXX | 0xD00-0xD7F | Standard read-only |
| 11 | 01 | 10XX | 0xD80-0xDBF | Standard read-only |
| 11 | 01 | 11XX | 0xDC0-0xDFF | Custom read-only |

⋮

Figure 1: CSR groups with address allocation [Waterman et al., 2016a].

CSRs Address Allocation Cont'd

⋮

| Hypervisor CSRs | | | | |
|-----------------|----|------|--------------|--------------------------------|
| 00 | 10 | XXXX | 0x200-0x2FF | Standard read/write |
| 01 | 10 | 0XXX | 0x600-0x67F | Standard read/write |
| 01 | 10 | 10XX | 0x680-0x6BF | Standard read/write |
| 01 | 10 | 11XX | 0x6C0-0x6FF | Custom read/write |
| 10 | 10 | 0XXX | 0xA00-0xA7F | Standard read/write |
| 10 | 10 | 10XX | 0xA80-0xABF | Standard read/write |
| 10 | 10 | 11XX | 0xAC0-0xAFF | Custom read/write |
| 11 | 10 | 0XXX | 0xE00-0xE7F | Standard read-only |
| 11 | 10 | 10XX | 0xE80-0xEBF | Standard read-only |
| 11 | 10 | 11XX | 0xEC0-0xEFF | Custom read-only |
| Machine CSRs | | | | |
| 00 | 11 | XXXX | 0x300-0x3FF | Standard read/write |
| 01 | 11 | 0XXX | 0x700-0x77F | Standard read/write |
| 01 | 11 | 100X | 0x780-0x79F | Standard read/write |
| 01 | 11 | 1010 | 0x7A0-0x7AF | Standard read/write debug CSRs |
| 01 | 11 | 1011 | 0x7B0-0x7BF | Debug-mode-only CSRs |
| 01 | 11 | 11XX | 0x7C0-0x7FF | Custom read/write |
| 10 | 11 | 0XXX | 0xB00-0xB7F | Standard read/write |
| 10 | 11 | 10XX | 0xB80-0xBBF | Standard read/write |
| 10 | 11 | 11XX | 0xBC0-0xBFF | Custom read/write |
| 11 | 11 | 0XXX | 0xF00-0xF7F | Standard read-only |
| 11 | 11 | 10XX | 0xF80-0xFB7F | Standard read-only |
| 11 | 11 | 11XX | 0xFC0-0xFFF | Custom read-only |

Figure 2: CSR groups with address allocation [Waterman et al., 2016a].

Machine Level CSRs

| Number | Privilege | Name | Description |
|-------------------------------|-----------|------------------|--|
| Machine Information Registers | | | |
| 0xF11 | MRO | mvendorid | Vendor ID. |
| 0xF12 | MRO | marchid | Architecture ID. |
| 0xF13 | MRO | mimpid | Implementation ID. |
| 0xF14 | MRO | mhartid | Hardware thread ID. |
| Machine Trap Setup | | | |
| 0x300 | MRW | mstatus | Machine status register. ie, ip (globally) |
| 0x301 | MRW | misa | ISA and extensions |
| 0x302 | MRW | medeleg | Machine exception delegation register. |
| 0x303 | MRW | mideleg | Machine interrupt delegation register. |
| 0x304 | MRW | mie | Machine interrupt-enable register. ie (locally) |
| 0x305 | MRW | mtvec | Machine trap-handler base address. ie |
| 0x306 | MRW | mcouneren | Machine counter enable. |
| Machine Trap Handling | | | |
| 0x340 | MRW | mscratch | Scratch register for machine trap handlers. |
| 0x341 | MRW | mepc | Machine exception program counter. where pc will go when an interrupt comes |
| 0x342 | MRW | mcause | Machine trap cause. ie |
| 0x343 | MRW | mtval | Machine bad address or instruction. |
| 0x344 | MRW | mip | Machine interrupt pending. ip (locally) |
| Machine Memory Protection | | | |
| 0x3A0 | MRW | pmpcfg0 | Physical memory protection configuration. |
| 0x3A1 | MRW | pmpcfg1 | Physical memory protection configuration, RV32 only. |
| 0x3A2 | MRW | pmpcfg2 | Physical memory protection configuration. |
| 0x3A3 | MRW | pmpcfg3 | Physical memory protection configuration, RV32 only. |
| 0x3B0 | MRW | pmpaddr0 | Physical memory protection address register. |
| 0x3B1 | MRW | pmpaddr1 | Physical memory protection address register. |
| | | ⋮ | |
| 0x3BF | MRW | pmpaddr15 | Physical memory protection address register. |

Figure 3: Machine level CSRs [Waterman et al., 2016a].

Machine Level CSRs Cont'd

| Number | Privilege | Name | Description |
|--|-----------|-----------------------------|---|
| Machine Counter/Timers | | | |
| 0xB00 | MRW | <code>mcycle</code> | Machine cycle counter. |
| 0xB02 | MRW | <code>minstret</code> | Machine instructions-retired counter. |
| 0xB03 | MRW | <code>mhpmcounter3</code> | Machine performance-monitoring counter. |
| 0xB04 | MRW | <code>mhpmcounter4</code> | Machine performance-monitoring counter. |
| | | ⋮ | |
| 0xB1F | MRW | <code>mhpmcounter31</code> | Machine performance-monitoring counter. |
| 0xB80 | MRW | <code>mcycleh</code> | Upper 32 bits of <code>mcycle</code> , RV32I only. |
| 0xB82 | MRW | <code>minstreth</code> | Upper 32 bits of <code>minstret</code> , RV32I only. |
| 0xB83 | MRW | <code>mhpmcounter3h</code> | Upper 32 bits of <code>mhpmcounter3</code> , RV32I only. |
| 0xB84 | MRW | <code>mhpmcounter4h</code> | Upper 32 bits of <code>mhpmcounter4</code> , RV32I only. |
| | | ⋮ | |
| 0xB9F | MRW | <code>mhpmcounter31h</code> | Upper 32 bits of <code>mhpmcounter31</code> , RV32I only. |
| Machine Counter Setup | | | |
| 0x320 | MRW | <code>mcountinhibit</code> | Machine counter-inhibit register. |
| 0x323 | MRW | <code>mhpmevent3</code> | Machine performance-monitoring event selector. |
| 0x324 | MRW | <code>mhpmevent4</code> | Machine performance-monitoring event selector. |
| | | ⋮ | |
| 0x33F | MRW | <code>mhpmevent31</code> | Machine performance-monitoring event selector. |
| Debug/Trace Registers (shared with Debug Mode) | | | |
| 0x7A0 | MRW | <code>tselect</code> | Debug/Trace trigger register select. |
| 0x7A1 | MRW | <code>tdata1</code> | First Debug/Trace trigger data register. |
| 0x7A2 | MRW | <code>tdata2</code> | Second Debug/Trace trigger data register. |
| 0x7A3 | MRW | <code>tdata3</code> | Third Debug/Trace trigger data register. |
| Debug Mode Registers | | | |
| 0x7B0 | DRW | <code>dcsr</code> | Debug control and status register. |
| 0x7B1 | DRW | <code>dpc</code> | Debug PC. |
| 0x7B2 | DRW | <code>dscratch0</code> | Debug scratch register 0. |
| 0x7B3 | DRW | <code>dscratch1</code> | Debug scratch register 1. |

Figure 4: Machine level CSRs [Waterman et al., 2016a].

Trap/Interrupt Setup CSRs

Table 6: User, supervisor and machine mode trap/interrupt setup CSRs.

| U Mode | S Mode | M Mode | Description |
|---------|------------|------------|--------------------------------|
| ustatus | sstatus | mstatus | Status register. |
| uie | sie | mie | Interrupt-enable register. |
| utvec | stvec | mtvec | Trap-handler base address. |
| | sedeleg | medeleg | Exception delegation register. |
| | sideleg | mideleg | Interrupt delegation register. |
| | scounteren | mcounteren | Counter enable. |
| | | misa | ISA and extensions. |


Trap Handling CSRs

Table 7: User, supervisor and machine mode trap handling CSRs.

| U Mode | S Mode | M Mode | Description |
|----------|----------|----------|-------------------------------------|
| uscratch | sscratch | mscratch | Scratch register for trap handlers. |
| uepc | sepc | mepc | Exception program counter. |
| ucause | scause | mcause | Trap cause. |
| utval | stval | mtval | Bad address or instruction. |
| uip | sip | mip | Interrupt pending. |

Trap/Interrupt Setup CSRs: `mstatus`

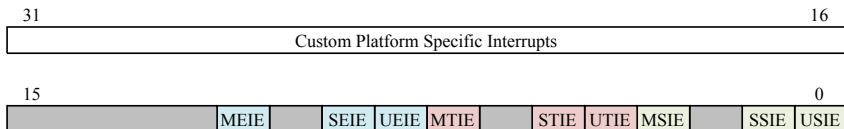
Table 8: Machine status register bit field descriptions.



| Bit field | Description |
|------------------|---|
| MIE, SIE, UIE | Interrupt-enable bits for each privilege mode. MIE for machine mode, SIE for supervisor mode, and UIE for user mode. |
| MPIE, SPIE, UPIE | x PIE holds the value of the interrupt-enable bit active prior to the trap, where $x \in \{M, S, U\}$. |
| MPP, SPP | x PP holds the previous privilege mode prior to the trap and can only hold privilege modes up to x . So MPP is two bits wide, SPP is one bit wide, and UPP is implicitly of bit width zero. |

Trap/Interrupt Setup CSRs: mie

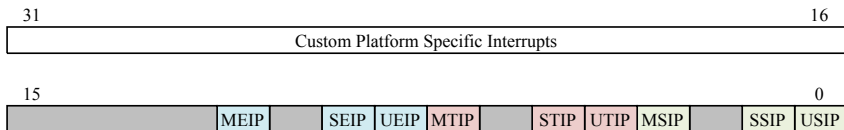
Table 9: Machine interrupt enable register bit field descriptions.



| Bit field | Description |
|------------------|---|
| MSIE, SSIE, USIE | These fields enable software interrupts for M-mode, S-mode and U-mode, respectively. |
| MTIE, STIE, UTIE | These fields enable timer interrupts for M-mode, S-mode, and U-mode, respectively. |
| MEIE, SEIE, UEIE | These fields enable external interrupts for M-mode, S-mode, and U-mode, respectively. |

Trap/Interrupt Handling CSRs: mip

Table 10: Machine interrupt pending register bit field descriptions.



| Bit field | Description |
|------------------|--|
| MSIP, SSIP, USIP | These fields correspond to software pending interrupts for M-mode, S-mode and U-mode, respectively. |
| MTIP, STIP, UTIP | These fields correspond to timer pending interrupts for M-mode, S-mode, and U-mode, respectively. |
| MEIP, SEIP, UEIP | These fields correspond to external pending interrupts for M-mode, S-mode, and U-mode, respectively. |

Trap/Interrupt Setup CSRs: `mtvec`


Table 11: Machine trap-handler base address register bit field descriptions.

| | | | |
|------------|--|---|------|
| 31 | | 2 | 0 |
| BASE[31:2] | | | MODE |

| Bit field | Description |
|------------|---|
| MODE | <p>This two bit field can be configured to following possible values:</p> <ul style="list-style-type: none"> • A value of 0 corresponds to Direct mode. All exceptions set pc to BASE field. • A value of 1 corresponds to Vectored mode. All exceptions set pc to (BASE + cause << 2). • Reserved for other values. |
| BASE[31:2] | <p>This field corresponds to 30 (most significant) bits of the 32-bit (word aligned) BASE address.</p> |

Trap/Interrupt Handling CSRs: mcause

Table 12: Machine trap cause register bit field descriptions.



| Bit field | Description |
|----------------|--|
| INTR | The INTR (Interrupt) bit is set if the trap was caused by an interrupt and is cleared in case of other traps (e.g. address fault). |
| Exception Code | This field contains a code identifying the last exception. |

Trap/Interrupt Handling CSRs: Exception Codes

Table 13: Machine cause register bit field values after trap.

| INTR | Exception Code | Description |
|------|----------------|--------------------------------|
| 1 | 0 | User software interrupt |
| 1 | 1 | Supervisor software interrupt |
| 1 | 2 | Reserved |
| 1 | 3 | Machine software interrupt |
| 1 | 4 | User timer interrupt |
| 1 | 5 | Supervisor timer interrupt |
| 1 | 6 | Reserved |
| 1 | 7 | Machine timer interrupt |
| ⋮ | ⋮ | ⋮ |
| 0 | 0 | Instruction address misaligned |
| 0 | 1 | Instruction access fault |
| 0 | 2 | Illegal instruction |
| 0 | 3 | Breakpoint |
| 0 | 4 | Load address misaligned |
| 0 | 5 | Load access fault |
| ⋮ | ⋮ | ⋮ |

Exceptions and Interrupts

- **Exception:** An *internal* event caused by program execution fault (e.g., page fault, misaligned memory access, arithmetic over-/under-flow)
- **Interrupt:** An *external* event caused outside of program execution (e.g. interrupt from a peripheral device UART, timer, Network Interface etc.)
- **Trap:** It is the forced transfer of execution control to supervisor/service-routine caused by exception or interrupt (not all exceptions/interrupts cause traps)

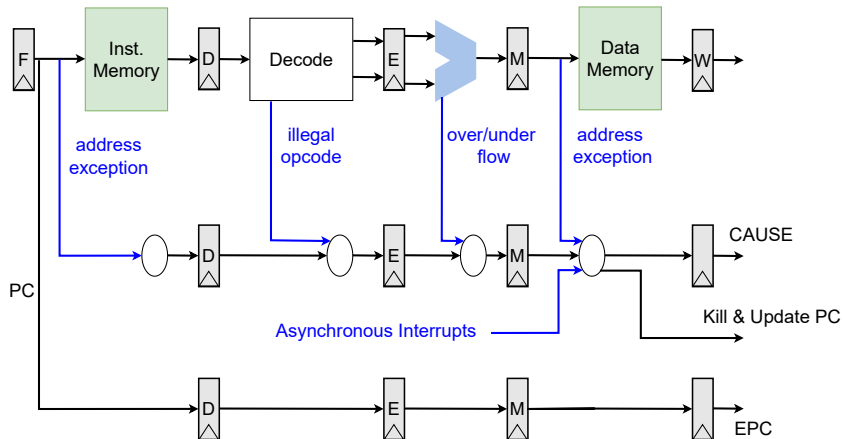
Exception and Interrupt Handling

- Exceptions are *synchronous* events
- Interrupts are *asynchronous* events
- How and where to handle multiple exceptions at different pipeline stages?
- How and where to handle asynchronous external interrupts

Exception and Interrupt Handling

- Exceptions are *synchronous* events
- Interrupts are *asynchronous* events
- How and where to handle multiple exceptions at different pipeline stages?
- How and where to handle asynchronous external interrupts
- Exceptions and interrupts are handled synchronously using **trap handlers**
- Exceptions and interrupts are handled at **memory** (or commit) stage

Exception and Interrupt Handling Cont'd



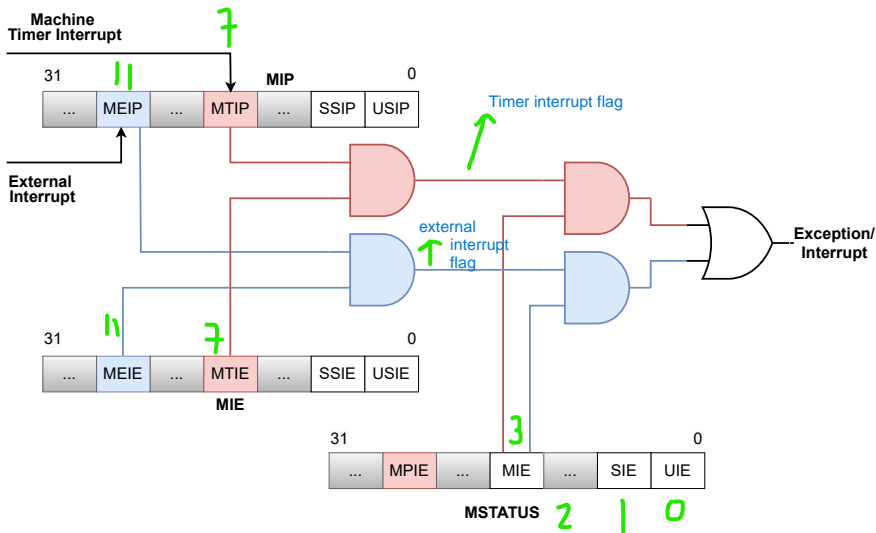
Exception and Interrupt Handling Cont'd

- Exception flags are held in pipeline till memory stage (or commit point)
- Interrupts (external events) are injected directly, overriding others, at memory stage
- At memory stage update **Cause** and **EPC** registers, kill all stages, update PC with trap handler address

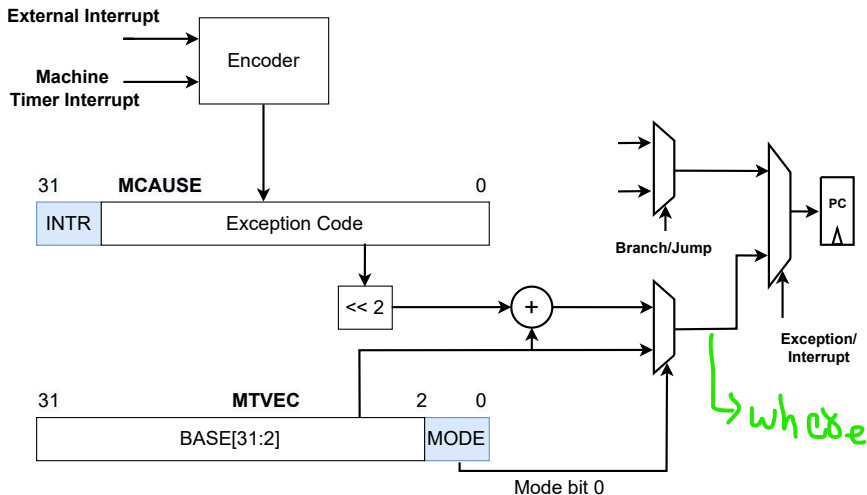
Trap Handler

- A Trap Handler at the start:
 - Saves **EPC** before enabling interrupts to allow nested interrupts
 - Reads the **Cause** register to find the source of exception/interrupt
- A Trap Handler at the end:
 - Uses a special instruction **MRET/SRET** to resume normal execution
 - Enables the interrupts
 - Restores the processor status

Trap Handling: Configuration



Trap Handling: Runtime



Example Startup File

```
.equ CSR_MSTATUS, 0x300
.equ MSTATUS_MIE, 0x00000008
.equ CSR_MTVEC, 0x305

# Main interrupt vector table entries
.global vtable
.type vtable, %object
.section .text .vector_table, "a", %progbits

# this entry is to align reset_handler at address 0x04
.word 0x00000013
j reset_handler
.align 2
vtable:
j default_interrupt_handler
.word 0
.word 0
j msip_handler
.word 0
.word 0
```

Example Startup File Cont'd

```
.word    0
j        mtip_handler
.word    0
.word    0
.word    0
.word    0
.word    0
.word    0
.word    0
.word    0
j        user_handler
.word    0
.word    0
```

Example Startup File Cont'd

```
# Weak aliases to point each exception handler to the
# 'default_interrupt_handler', unless the application defines
# a function with the same name to override the reference.

.weak msip_handler
.set msip_handler, default_interrupt_handler
.weak mtip_handler
.set mtip_handler, default_interrupt_handler
.weak user_handler
.set user_handler, default_interrupt_handler

# Assembly 'reset handler' function to initialize core CPU registers
.
.section .text.default_interrupt_handler,"ax",%progbits

.global reset_handler
.type reset_handler,@function

reset_handler:
# Set mstatus bit MIE = 1 (enable M mode interrupts)
li      t0, 8
csrrs   zero, CSR_MSTATUS, t0
```


Example Startup File Cont'd

```
# Load the initial stack pointer value.
    la    sp, _sp

# Set the vector table's base address.
    la    a0, vtable
    addi  a0, a0, 1
    csrw  CSR_MTVEC, a0

# Call user 'main(0,0)' (.data/.bss sections initialized there)
    li    a0, 0
    li    a1, 0
    call  main

# A 'default' handler, in case an interrupt triggers without its
  handler defined
default_interrupt_handler:
    j     default_interrupt_handler
```

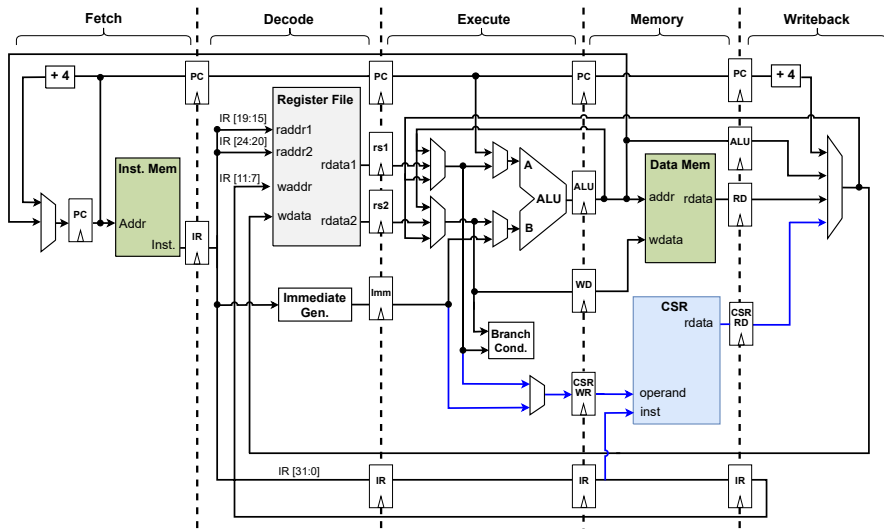
Example Interrupt Service Routine

```
# RISC-V Interrupt Service Routines (ISRs)
# ALL supported ISRs should be put here

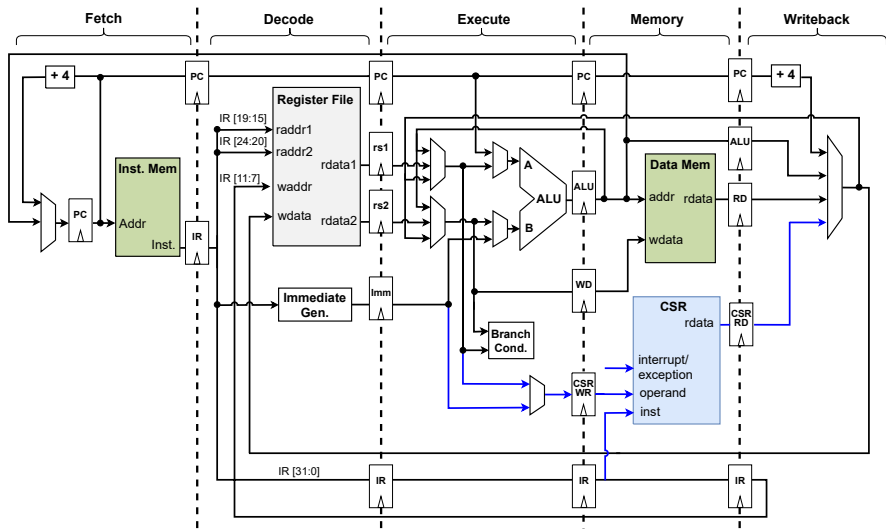
.section .text.isr

# User interrupt handler
.globl user_handler
user_handler:
    nop
    # you can call user ISR here and then return using 'mret'
    mret
```

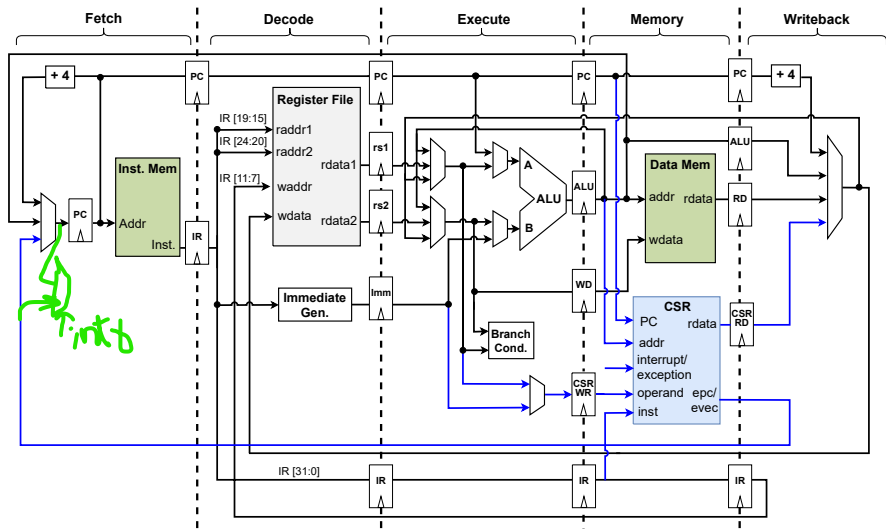
Pipelined Datapath with Exception/Interrupt Handling



Pipelined Datapath with Exception/Interrupt Handling



Pipelined Datapath with Exception/Interrupt Handling



csrwr x9, mcause, x10

Suggested Reading

- Read User Manual for the instruction set and its architecture [[Waterman et al., 2016b](#)].
- For Control and Status register description consult Privileged architecture manual [[Waterman et al., 2016a](#)].
- Read Section 5.14 of [[Patterson and Hennessy, 2021](#)].

Acknowledgment

- Preparation of this material was partly supported by Lampro Mellon Pakistan.

References



Patterson, D. and Hennessy, J. (2021).

Computer Organization and Design RISC-V Edition: The Hardware Software Interface, 2nd Edition.

Morgan Kaufmann.



Waterman, A., Lee, Y., Avizienis, R., Patterson, D. A., and Asanovic, K. (2016a).

The risc-v instruction set manual volume ii: Privileged architecture version 1.9.

EECS Department, University of California, Berkeley, Tech. Rep. UCB/EECS-2016-129.



Waterman, A., Lee, Y., Patterson, D. A., and Asanović, K. (2016b).

The risc-v instruction set manual, volume i: User-level isa, version 2.1.

EECS Department, University of California, Berkeley, Tech. Rep. UCB/EECS-2016-129.